

RL using Temporal Difference

By Sefunmi Ashiru

Environment: OpenAI gym

FrozenLake8x8-v0


(S: starting point, safe)

(F: frozen surface, safe)

(H: hole, fall to your doom)

(G: goal, where the frisbee is located)

(■: agent)



Q-Learning

Tabular method

Problems in which the state and actions spaces are small enough for approximate value functions to be represented as arrays and tables.

	Action 1 (up)	Action 2 (left)	Action 3 (right)	Action 4 (down)
State 1	$Q(s_1, a_1)$	$Q(s_1, a_2)$	$Q(s_1, a_3)$	$Q(s_1, a_4)$
State 2	$Q(s_2, a_1)$	$Q(s_2, a_2)$	$Q(s_2, a_3)$	$Q(s_2, a_4)$
State 3	$Q(s_3, a_1)$	$Q(s_3, a_2)$	$Q(s_3, a_3)$	$Q(s_3, a_4)$
State k	$Q(s_k, a_1)$	$Q(s_k, a_2)$	$Q(s_k, a_3)$	$Q(s_k, a_4)$

Pros

- Doesn't require the transition probability matrix like in Dynamic Programming.
- Updates Q/"state action pair" values incrementally. Don't need to wait till episode ends.

Cons

- Q-Learning can behave poorly in some stochastic environments.

Algorithm

Episodes = 200000

Learning Rate:

$\alpha = .648$

Discounted Rewards:

$\gamma = .9$

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

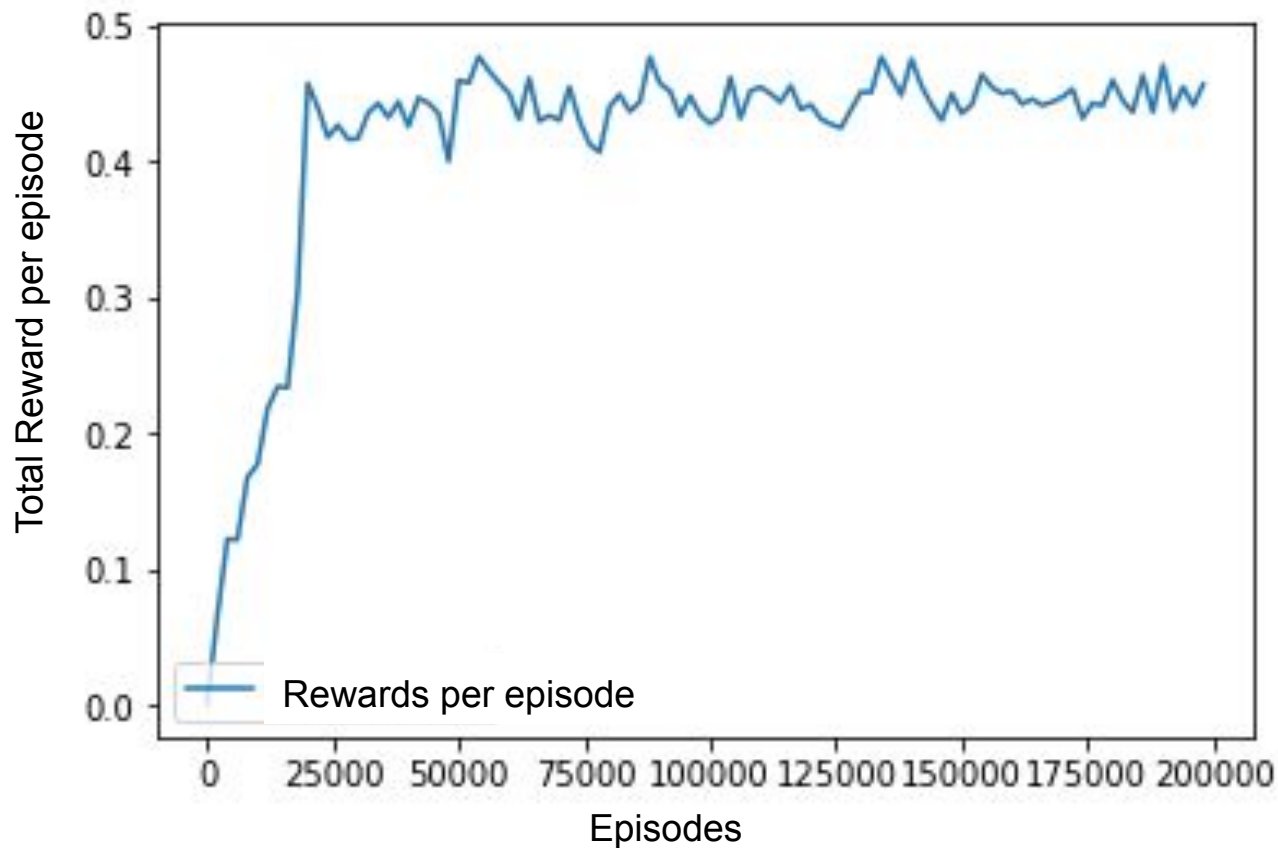
$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

[1] (Salloum, 2018)

Rewards per episode



Bibliography

1. Salloum, Z. (2018, December 16). Summary of Tabular Methods in Reinforcement Learning [Web log post]. Retrieved 2020, from <https://towardsdatascience.com/summary-of-tabular-methods-in-reinforcement-learning-39d653e904af#:~:text=Introduction,represented%20as%20arrays%20and%20tables>.

